David Tianhao Li

🏠 tianhao.li ✉ tianhao.li@duke.edu in LinkedIn G Google Scholar ⌂ Github

Overview

I have spent nearly 3 years working in AI safety and security, 1.5 years in industry and 1.5 years in academia, with a recent focus on building safe and secure autonomous agents and multi-agent systems. My long-term goal is to build trustworthy and responsible intelligence systems for high-stakes domains (incl. healthcare, finance, science, etc.) with my interdisciplinary background (MSc: healthcare/physics; BEng: computer science/information security) and work experiences (ByteDance: AI agent; TopSec/NSFocus: offensive AI security).

Education

Duke University   **2024/08 - 2026/05 (Expected)**
Master of Science Candidate in DKU Medical Physics   *GPA: 3.5/4*
Research: Self-Evolving Healthcare AI
With: Prof. Neil Gong (ECE/CS, Duke University), Prof. Chaowei Xiao (ECE/CS, John Hopkins University)

North China University of Technology   **2020/09 - 2024/06**
Bachelor of Engineering in Information Security   *GPA: 3.9/4*
Project: A Multi-modal LLM Red Teaming Platform Based on NVIDIA/Garak
Honors: Outstanding Thesis Award (1/101), Outstanding Student Scholarship (1/329)

Employments

ByteDance   **2025/11 – Present**
Agent Research (Intern) - Trajectory & Self-Evolving - with FTE opportunity   *Onsite @ Beijing, China*

- Led a project on LLM agent trajectory (a.k.a. trace or transcript) evaluation and alignment.
- Led a project focused on developing a self-evolving ToB agentic capability for Seed Foundation Model, utilizing bad-case analysis and synthesis-based automated data manufacturing pipeline for reinforcement learning (RL).

TOPSEC   **2024/03 – 2024/09**
Security Researcher (Intern) - Adversarial ML - with FTE return offer (declined)   *Onsite @ Beijing, China*

- Developed a prototype software system for ML model adversarial robustness evaluation;
- Contributed to a phishing email detection system using SFT-LLMs, with responsibility for dataset collection and refinement.

NSFOCUS   **2023/09 – 2024/03**
Security Researcher (Intern) - LLM Red Teaming   *Onsite @ Beijing, China*

- Contribute to algorithms design and prototype development of NSFOCUS LLM Security Assessment System (LSAS);
- Performed offensive security test on multiple public and private large language models to evaluate vulnerabilities;
- Reviewed academic papers and Gartner reports on trustworthy AI to enrich the threat intelligence database.

Projects

AI Safety Student Team, Havard University   **2026/2 – 2026/4 (Expected)**
AISST Technical Fellowship (Part-time)   *Remote @ MA, United States*

- Complete a 8 week curriculum with hands-on work on reward misspecification, RLHF, goal misgeneralization, and red teaming, with a focus on scalable oversight, adversarial attacks, and interpretability in AI systems.

Parker & Lawrence Research   **2026/2 – Present**
Founding Member of Technical Staff (Contract, Part-time)   *Remote @ London, United Kingdom*

- Build AI Quantified (AIQ), a LangGraph-based multi-agent system for AI technology intelligence quantitative analysis.

SIGMIR   **2025/01 – Present**
Co-founder & CEO   *Remote @ CA, United States*

- Founded SIGMIR, a 501(c)(3) non-profit; built the core team, led early operations, and oversee all projects.
- Facilitating connections between academia and industry and among individuals from diverse disciplines to cultivate collaborative research and innovation.

Funding

Publications

- **Tianhao Li**, Chuangxin Chu, Yujia Zheng, Bohan Zhang, Neil Zhenqiang Gong, and Chaowei Xiao. (2025). A2ASecBench: A Protocol-Aware Security Benchmark for Agent-to-Agent Multi-Agent Systems. *ICLR 2026 Accepted* - *openreview.net/forum?id=LfdFnakqGJ*.

- **Tianhao Li**, Jingyu Lu, Chuangxin Chu, Tianyu Zeng, Yujia Zheng, Mei Li, Haotian Huang, Bin Wu, Zuoxian Liu, Kai Ma, Xuejing Yuan, Xingkai Wang, Keyan Ding, Huajun Chen, and Qiang Zhang. (2024). SciSafeEval: A comprehensive benchmark for safety alignment of large language models in scientific tasks. *AAAI 2025* *AI for Cybersecurity - ACM TIST Under Review* - *arXiv:2410.03769*.

- **Tianhao Li**, Jie Liu, Wenbo Shi, Yiyun Lu, Leran Chen, Yongshuai Li. (2026). A Survey of Trajectory Evaluation and Alignment for LLM Agents. ARR Jan 2026 Submission.

- Yujia Zheng[1st], **Tianhao Li**[1st], Haotian Huang, Tianyu Zeng, Jingyu Lu, Chuangxin Chu, Yuekai Huang, Ziyou Jiang, Qian Xiong, Yuyao Ge and Mingyang Li (2025). Are all prompt components value-neutral? Understanding the heterogeneous adversarial robustness of dissected prompt in large language models. *EACL 2026 Oral* - *arXiv:2508.01554*.

- Tianyu Zeng[1st], **Tianhao Li**[1st], Yujia Zheng[1st], Kaizhu Huang, and Zhenyu Yang. (2025). Do Synthetic Medical Images Pose Unseen Risks to Clinical Segmentation Tasks?. *ACM TIST Under Review*.

- Weizhi Ma, Ying Li, **Tianhao Li**, Haowei Yang, Zhengping Li, Lijun Wang, and Junyu Xuan. (2025). SFSWTS: A spatial-frequency shifted windows and time self-attention network for EEG emotion recognition. *Neurocomputing* - *doi.org/10.1016/j.neucom.2025.130309*.

- Qian Xiong, Yuekai Huang, Ziyou Jiang, Zhiyuan Chang, Yujia Zheng, **Tianhao Li**, and Mingyang Li. (2025). Butterfly effects in toolchains: A comprehensive analysis of failed parameter filling in LLM tool-agent systems. *EMNLP 2025 Findings* - *arXiv:2507.15296*.

- Qian Xiong, Yuekai Huang, Yujia Zheng, **Tianhao Li**, Ziyou Jiang, Zhiyuan Chang, Zhaoyang Li, Huanxiang Feng, Mingyang Li. (2026). Retrogradus: Addressing Intent Deviation in Tool-Using Agents via Reverse Data Synthesis and Mutation on Critical Parameters. *IJCAI 2026 Submission* - *arXiv:2601.15120*.

- Qiang Zhang, Xiang Zhuang, Chenyi Zhou, Yihang Zhu, Tong Xu, **Tianhao Li**, Dianbo Liu, Shengchao Liu, Keyan Ding, Emine Yilmaz, Haofen Wang, and Huajun Chen. (2025). Beyond Accuracy: Comprehensive Alignment for AI-Driven Molecular Design. *Nature Computational Science* *Major Revision*.

- Yuanyuan Wei, Xianxian Liu, Yao Mu, Changran Xu, Guoxun Zhang, **Tianhao Li**, Zida Li, Wu Yuan, Ho-Pui Ho, and Mingkun Xu. (2025). From droplets to diagnosis: AI-driven imaging and system integration in digital nucleic acid amplification testing. *Biosensors and Bioelectronics* - *doi.org/10.1016/j.bios.2025.117741*.

- Guofan Zhang, Xuanzhi Wang, Qirui Sun, Hanchi Zhao, Jiaao Han, Haoyuan Yang, Hanyu Wang, Chengshu Tian, Aizhen Kuang, Lanxiang Lai, Chulong Zhang, **Tianhao Li** and Rui Liu (2025). A versatile method for accurately predicting electronic absorption spectra of tetrapyrrole macrocycles. *ChemRxiv* - *doi.org/10.26434/chemrxiv-2025-9gsp2*.

- Yepeng Feng, Yujia Zheng., **Tianhao Li** and Xuan Tian. (2025). Personalized Query Recommendation for Large Language Models Guided by Dynamic User Profiles and Bidirectional Intention. *Neurocomputing* *Submitted* - dx.doi.org/10.2139/ssrn.6232769.

- Yujia Zheng, Tianyu Zeng, Haotian Huang, Chulong Zhang, **Tianhao Li**, Jingyu Lu, Chuangxin Chu, Manju Liu, Chunhao Wang, Kaizhu Huang, Fang-Fang Yin and Zhenyu Yang. (2025). Assessing the Risks of Synthetic Data in Deep Learning-Based PET Tumor Segmentation. *International Journal of Radiation Oncology, Biology, Physics* - *doi.org/10.1016/j.ijrobp.2025.06.3861*

- Weizhi Ma, Yujia Zheng, **Tianhao Li**, Zhengping Li, Ying Li, and Lijun Wang. (2024). A comprehensive review of deep learning in EEG-based emotion recognition: classifications, trends, and practical implications. *PeerJ Computer Science* - *doi.org/10.7717/peerj-cs.2065*.

- Yuyao Ge, Lingrui Mei, Zenghao Duan, **Tianhao Li**, Yujia Zheng, Yiwei Wang, Lexin Wang, Jiayu Yao, Tianyu Liu, Yujun Cai, Baolong Bi, Fangda Guo, Jiafeng Guo, Shenghua Liu, and Xueqi Cheng. (2025). A Survey of Vibe Coding with Large Language Models. *arXiv:2510.12399*.

- Gregor von Laszewski, Wesley Brewer, Jeyan Thiyagalingam, Juri Papay, Armstrong Foundjem, Piotr Luszczek, Murali Emani, Shirley V. Moore, Vijay Janapa Reddi, Matthew D. Sinclair, Sebastian Lobentanzer, Sujata Goswami, Benjamin Hawks, Marco Colombo, Nhan Tran, Christine R. Kirkpatrick, Abdulkareem Alsudais, Gregg Barrett, **Tianhao Li**, Kirsten Morehouse, Shivaram Venkataraman, Rutwik Jain, Kartik Mathur, Victor Lu, Tejinder Singh, Khojasteh Z. Mirza, Kongtao Chen, Sasidhar Kunapuli, Gavin Farrell, Renato Umeton, and Geoffrey C. Fox. (2025). AI Benchmarks Carpentry and Democratization. *Frontiers in High Performance Computing, Section Benchmarking* *Under Review - arXiv:2512.11588*.

Technical Reports

- Ghosh et al. (2025). AILuminate: Introducing v1.0 of the AI Risk and Reliability Benchmark from MLCommons. *MLCommons - arXiv:2503.05731* - mlcommons.org/en/ailuminate.

- Goel et al. (2025). AILuminate Security: Introducing v0.5 of the Jailbreak Benchmark from MLCommons. *MLCommons - wp-content* - mlcommons.org/ailuminate/jailbreak/.

- McGregor et al. (2025). Agentic Product Maturity Ladder V0.1. *MLCommons - wp-content* - mlcommons.org/ailuminate/agentic/.

- Maple et al. (2026). A Robust, Defensible, and Reproducible Methodology for Benchmarking Single-Turn Jailbreak Attacks on Large Language Models. *MLCommons - wp-content* - mlcommons.org/2026/02/jailbreak-0-7/.

Academic Talks

- Poster Presentation - A2ASecBench: A Protocol-Aware Security Benchmark for Agent-to-Agent Multi-Agent Systems. *The Fourteenth International Conference on Learning Representations (ICLR 2026)*, Apr 2026, Rio de Janeiro, Brazil.

- Oral Presentation - Are All Prompt Components Value-Neutral? Understanding the Heterogeneous Adversarial Robustness of Dissected Prompt in Large Language Models (15 mins). *The 19th Conference of the European Chapter of the Association for Computational Linguistics (EACL 2026)*, Mar 2026, Rabat, Morocco.

- Poster Presentation - Butterfly Effects in Toolchains: A Comprehensive Analysis of Failed Parameter Filling in LLM Tool-Agent Systems *The 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP 2025)*, Nov 2025, Jiangsu, China.

- Invited talk - Toward Trustworthy Generative Foundation Models and Autonomous Agent System (25 mins). *Institute of Software, Chinese Academy of Sciences*, Aug 2025, Beijing, China.

- Contributed talk - Guardians of Trust: Safety & Privacy for Healthcare Foundation Models (15 mins). *Duke Kunshan Graduate Student Colloquium*, Apr 2025, Jiangsu, China.

- Oral Presentation - SciSafeEval: A Comprehensive Benchmark for Safety Alignment of Large Language Models in Scientific Tasks (20 mins). *The 39th Annual AAAI Conference on Artificial Intelligence (AAAI-2025), Artificial Intelligence for Cyber Security (AICS)*, Mar 2025, Philadelphia, PA, USA.

Professional Services

- Journal Referee: ACM TIST, IEEE TAI, IEEE TBME, IEEE JBHI, IJHCI, ESWA, EAAI, RESS, COMPUT NETW, PATTERN RECOGN LETT., PATTERN RECOGNIT.

- Program Committee Member: ICLR, AAAI, ACL ARR, IJCAI, COLING, IJCNN, ISBI, CHIL, LLMSEC

Mentoring Experience

- Haotian Huang, NCUT UG'27, China

- Akshath Mangudi, VIT-AP University UG'26, India

Open Source Activities

NVIDIA/garak
github.com/nvidia/garak - Generative AI red-teaming & assessment kit   *Contributor*
- Extended from single modality (text-to-text) only to multimodal (any-to-any) capability with #587: implemented a FigStep-based [(image+text)-to-text] jailbreaking pipeline for LLaVA - the first multi-modal red-teaming tool.
- Improved system reliability through bug fixes (#296, #401) and enhanced exception handling (#566) across core modules.

Google/adk-samples
github.com/google/adk-samples - A collection provides ready-to-use agents built on top of the google-adk   *Contributor*
- Enhancing Data Science Agent's configurability by introducing environment-based dataset management (#220, #221).

A2AProject/a2a-python
github.com/a2aproject/a2a-python - Official python SDK for the Google Agent2Agent (A2A) protocol   *Contributor*
- Enhanced reliability by redesigning the streaming error-handling architecture to accurately surface server status and messages. (#502, #505)

SIGMIR-ORG/TARMAS
github.com/sigmir-org/tarmas - A framework for threat analysis and Risk metrics in agent systems   *Owner*
- Lead project development and research direction, overseeing design, implementation, and maintenance

Media Coverage

- Concordia AI. AI Action Summit: Public Interest AI. https://concordia-ai.com/wp-content/uploads/2024/11/Concordia-AI-French-AI-Action-Summit-Public-Interest-AI-contribution.pdf
- Concordia AI. AI Action Summit: AI of Trust - Security & Safety. https://concordia-ai.com/wp-content/uploads/2024/11/Concordia-AI-French-AI-Action-Summit-AI-of-Trust-Security-Safety-contribution.pdf
- Shanghai AI Lab. Frontier AI Risk Management Framework. https://research.ai45.shlab.org.cn/safework-f1-framework.EN.pdf
- IEEE Spectrum. One Chatbot Safety Benchmark To Test Them All: AILuminate aims to create an industry standard for large language model safety. https://spectrum.ieee.org/ai-safety-standard
- BusinessWire. MLCommons Launches AILuminate, First-of-Its-Kind Benchmark to Measure the Safety of Large Language Models. https://businesswire.com/news/home/20241204285618/en/MLCommons-Launches-AILuminate-First-of-Its-Kind-Benchmark-to-Measure-the-Safety-of-Large-Language-Models
- Binary Verse AI. Vibe Coding Tools: The Definitive 2025 Guide, Based on a Landmark AI Survey. https://binaryverseai.com/best-vibe-coding-tools-top-agentic-cursor-ai/
- PYMNTS Logo. Vibe Coding Won't Replace Humans Anytime Soon, Data Shows https://www.pymnts.com/artificial-intelligence-2/2025/vibe-coding-wont-replace-humans-anytime-soon-data-shows/

Competition Awards

- 2024/08, Second Prize in 17th CISCN Security Project Contest   *National*
- 2024/07, Top 5 out of 89 in Galaxy Generative AI Safety Contest Attack Track   *National*
- 2024/02, S Prize in 2024 Mathematical Contest In Modeling (MCM)   *National*
- 2023/11, Second Prize in 11th Digital Media Technology and Creativity Contest   *National*
- 2023/11, Third Prize in 8th National Cryptography Contest   *National*
- 2023/08, Third Prize in 16th CISCN Security Project Contest   *National*
- 2023/06, Third Prize in 16th CISCN AWDP (Attack with Defense Plus) Contest   *Regional*
- 2023/04, First Prize in 13th MathorCup Mathematical Modeling Challenge   *National*
- 2023/04, First Prize in 14th Lanqiao Software Development and Algorithm Contest   *Provincial*
- 2022/06, Winning Prize in 15th CISCN Security Project Contest   *National*
- 2022/06, Third Prize in 15th CISCN CTF (Capture the Flag) Contest   *Regional*
- 2021/06, Third prize in 12th Lanqiao Software Development and Algorithm Contest   *National*
- 2021/04, First Prize in 12th Lanqiao Software Development and Algorithm Contest   *Provincial*